

# Lab 3 :

---

Author: Badr TAJINI - DevOps Data for SWE - ESIEE - 2024/2025

---

## Lab: How to Deploy Your Apps

---

We will work step-by-step on lab exercises to help you understand different orchestration methods for deploying applications, as covered in Chapter 3, "How to Deploy Your Apps." We will explore:

- **Server Orchestration** using Ansible
  - **VM Orchestration** using Packer and OpenTofu (an open-source Terraform fork)
  - **Container Orchestration** using Docker and Kubernetes
- 

## Table of Contents

---

- [Prerequisites](#)
- [Part 1: Server Orchestration with Ansible](#)
  - [Step 1: Set Up the Ansible Environment](#)
  - [Step 2: Creating EC2 Instances with Ansible](#)
  - [Step 3: Configuring Dynamic Inventory](#)
  - [Step 4: Deploying the Sample Node.js Application](#)
  - [Step 5: Setting Up Nginx as a Load Balancer](#)
  - [Step 6: Implementing Rolling Updates](#)
- [Part 2: VM Orchestration with Packer and OpenTofu](#)
  - [Step 1: Building a VM Image Using Packer](#)
  - [Step 2: Deploying the VM Image Using OpenTofu](#)
  - [Step 3: Deploying an Application Load Balancer \(ALB\)](#)
  - [Step 4: Implementing Rolling Updates with ASG Instance Refresh](#)
- [Part 3: Container Orchestration with Docker and Kubernetes](#)
  - [Step 1: Building and Running the Docker Image Locally](#)
  - [Step 2: Deploying the Application to a Local Kubernetes Cluster](#)
  - [Step 3: Performing a Rolling Update](#)

- [Step 4: Deploying a Kubernetes Cluster in AWS Using EKS](#)
- [Step 5: Pushing a Docker Image to Amazon ECR](#)
- [Step 6: Tag and Push the Docker Image to ECR](#)
- [Step 7: Deploying the Sample Application to the EKS Cluster](#)
- [Part 4: Deploying Applications Using Serverless Orchestration with AWS Lambda](#)
- [Cleanup](#)
- [Conclusion](#)

---

## Prerequisites

---

Before starting, ensure you have the following:

- **AWS Account:** An AWS account with permissions to create and manage resources.
- **AWS CLI Installed and Configured:** AWS CLI configured with your AWS credentials.
- **Ansible Installed:** Version 2.9 or later.
- **Packer Installed:** To build VM images.
- **OpenTofu Installed:** Install [OpenTofu](#).
- **Docker Desktop Installed:** Ensure Kubernetes is enabled.
- **Kubectl Installed:** The Kubernetes command-line tool.
- **Git Installed:** To clone repositories from GitHub.
- **SSH Key Pair:** To access EC2 instances via SSH.

---

## Part 1: Server Orchestration with Ansible

---

### Step 1: Set Up the Ansible Environment

#### 1.1 Clone the Sample Code Repository

```
cd devops-book/td3/scripts/ansible
```

#### 1.2 Install Required Ansible Collections

```
ansible-galaxy collection install amazon.aws
```

## 1.3 Configure AWS Credentials for Ansible

### Option 1: Use AWS CLI Configuration

```
aws configure
```

### Option 2: Set Environment Variables

```
export AWS_ACCESS_KEY_ID=your_access_key_id
export AWS_SECRET_ACCESS_KEY=your_secret_access_key
export AWS_DEFAULT_REGION=us-east-2
```

## Step 2: Creating EC2 Instances with Ansible

### 2.1 Create a Variables File

Create `sample-app-vars.yml` :

```
num_instances: 3
base_name: sample_app_instances
http_port: 8080
```

### 2.2 Run the Ansible Playbook to Create EC2 Instances

```
ansible-playbook -v create_ec2_instances_playbook.yml --extra-vars "@sample-app-
vars.yml"
```

### 2.3 Verify the EC2 Instances

- Navigate to the **EC2 Dashboard** in AWS Management Console.
- Confirm three instances are running with the tag `Ansible=sample_app_instances`.

## Step 3: Configuring Dynamic Inventory

### 3.1 Create Inventory Configuration

Create `inventory.aws_ec2.yml` :

```
plugin: amazon.aws.aws_ec2
regions:
  - us-east-2
keyed_groups:
  - key: tags.Ansible
leading_separator: ''
```

## 3.2 Create Group Variables

Create `group_vars/sample_app_instances.yml` :

```
ansible_user: ec2-user
ansible_ssh_private_key_file: ansible-ch3.key
ansible_host_key_checking: false
```

# Step 4: Deploying the Sample Node.js Application

## 4.1 Create the Application Deployment Playbook

Create `configure_sample_app_playbook.yml` :

```
- name: Configure servers to run the sample-app
  hosts: sample_app_instances
  gather_facts: true
  become: true
  roles:
    - role: nodejs-app
    - role: sample-app
      become_user: app-user
```

## 4.2 Create the `nodejs-app` Role

Create `roles/nodejs-app/tasks/main.yml` :

```

- name: Add Node.js packages to yum
  shell: curl -fsSL https://rpm.nodesource.com/setup_21.x | bash -

- name: Install Node.js
  yum:
    name: nodejs

- name: Create app user
  user:
    name: app-user

- name: Install PM2
  npm:
    name: pm2
    global: true

- name: Configure PM2 to run at startup as the app user
  shell: |
    su - app-user -c "pm2 startup systemd -u app-user --hp /home/app-user"
    systemctl enable pm2-app-user

```

### 4.3 Create the `sample-app` Role

Copy `app.js` and `app.config.js` into `roles/sample-app/files/`.

- `app.js` :

```

const http = require('http');
const server = http.createServer((req, res) => {
  res.end('Hello, World!\n');
});
server.listen(8080, () => {
  console.log('Listening on port 8080');
});

```

- `app.config.js` :

```

module.exports = {
  apps : [{
    name    : "sample-app",
    script  : "./app.js",
    exec_mode: "cluster",
    instances: "max",
    env: {
      "NODE_ENV": "production"
    }
  }]
}

```

Create `roles/sample-app/tasks/main.yml`:

```

- name: Copy sample app
  copy:
    src: "{{ item }}"
    dest: "/home/app-user/"
    owner: app-user
    group: app-user
    mode: 0644
  with_fileglob:
    - "files/*"

- name: Start sample app using PM2
  shell: pm2 start app.config.js
  args:
    chdir: /home/app-user/
  become_user: app-user

- name: Save PM2 process list
  shell: pm2 save
  become_user: app-user

```

## 4.4 Run the Application Deployment Playbook

```

ansible-playbook -v -i inventory.aws_ec2.yml configure_sample_app_playbook.yml

```

## 4.5 Verify the Application is Running

Get public IPs:

```
aws ec2 describe-instances \
  --filters "Name=tag:Ansible,Values=sample_app_instances" \
  --query "Reservations[*].Instances[*].PublicIpAddress" \
  --output text
```

Test the application:

```
curl http://<EC2_INSTANCE_PUBLIC_IP>:8080
```

## Step 5: Setting Up Nginx as a Load Balancer

### 5.1 Create a Variables File for Nginx

Create `nginx-vars.yml`:

```
num_instances: 1
base_name: nginx_instances
http_port: 80
```

### 5.2 Run the Playbook to Create an EC2 Instance for Nginx

```
ansible-playbook -v create_ec2_instances_playbook.yml --extra-vars "@nginx-vars.yml"
```

### 5.3 Create Group Variables for Nginx

Create `group_vars/nginx_instances.yml`:

```
ansible_user: ec2-user
ansible_ssh_private_key_file: ansible-ch3.key
ansible_host_key_checking: false
```

### 5.4 Create the Nginx Playbook

Create `configure_nginx_playbook.yml`:

```
- name: Configure servers to run Nginx
  hosts: nginx_instances
  gather_facts: true
  become: true
  roles:
    - role: nginx
```

## 5.5 Create the `nginx` Role

Create `roles/nginx/templates/nginx.conf.j2`:

```
user nginx;
worker_processes auto;
error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    log_format main '$remote_addr - $remote_user [$time_local] "$request" '
        '$status $body_bytes_sent "$http_referer" '
        '"$http_user_agent" "$http_x_forwarded_for"';

    access_log /var/log/nginx/access.log main;

    include /etc/nginx/mime.types;
    default_type application/octet-stream;

    upstream backend {
        {% for host in groups['sample_app_instances'] %}
        server {{ hostvars[host]['ansible_host'] }}:8080;
        {% endfor %}
    }

    server {
        listen 80;
        listen [::]:80;

        location / {
            proxy_pass http://backend;
        }
    }
}
```

Create `roles/nginx/tasks/main.yml`:



```

- name: Install Nginx
  yum:
    name: nginx
    state: present

- name: Copy Nginx config
  template:
    src: nginx.conf.j2
    dest: /etc/nginx/nginx.conf
  notify:
    - restart nginx

- name: Start and enable Nginx
  service:
    name: nginx
    state: started
    enabled: true

handlers:
  - name: restart nginx
    service:
      name: nginx
      state: restarted

```

## 5.6 Run the Nginx Playbook

```
ansible-playbook -v -i inventory.aws_ec2.yml configure_nginx_playbook.yml
```

## 5.7 Verify the Load Balancer

Get Nginx public IP:

```

aws ec2 describe-instances \
  --filters "Name=tag:Ansible,Values=nginx_instances" \
  --query "Reservations[*].Instances[*].PublicIpAddress" \
  --output text

```

Test the load balancer:

```
curl http://<NGINX_PUBLIC_IP>
```

## Step 6: Implementing Rolling Updates

## 6.1 Enable Rolling Updates in the Playbook

Modify `configure_sample_app_playbook.yml` :

```
- name: Configure servers to run the sample-app
  hosts: sample_app_instances
  gather_facts: true
  become: true
  serial: 1
  max_fail_percentage: 30
  roles:
    - role: nodejs-app
    - role: sample-app
      become_user: app-user
```

## 6.2 Update the Application

Modify `roles/sample-app/files/app.js` :

```
res.end('DevOps Base!\n');
```

## 6.3 Run the Application Deployment Playbook Again

```
ansible-playbook -v -i inventory.aws_ec2.yml configure_sample_app_playbook.yml
```

## 6.4 Verify the Rolling Update

Continuously test the application:

```
while true; do curl http://<NGINX_PUBLIC_IP>; sleep 1; done
```

---

# Part 2: VM Orchestration with Packer and OpenTofu

---

## Step 1: Building a VM Image Using Packer

### 1.1 Set Up the Working Directory

```
mkdir -p devops_base/td3/scripts/packer  
cd devops_base/td3/scripts/packer
```

## 1.2 Create the Packer Template

Create `sample-app.pkr.hcl` :

```

packer {
  required_plugins {
    amazon = {
      version = ">= 1.0.0"
      source  = "github.com/hashicorp/amazon"
    }
  }
}

variable "aws_region" {
  type    = string
  default = "us-east-2"
}

source "amazon-ebs" "amazon_linux" {
  ami_name      = "packer-sample-app-{{timestamp}}"
  instance_type = "t2.micro"
  region        = var.aws_region
  source_ami_filter {
    filters = {
      name                = "amzn2-ami-hvm-*-x86_64-gp2"
      root-device-type    = "ebs"
      virtualization-type = "hvm"
    }
    owners      = ["amazon"]
    most_recent = true
  }
  ssh_username = "ec2-user"
}

build {
  sources = ["source.amazon-ebs.amazon_linux"]

  provisioner "file" {
    sources      = ["app.js", "app.config.js"]
    destination = "/tmp/"
  }

  provisioner "shell" {
    inline = [
      "sudo yum update -y",
      "curl -fsSL https://rpm.nodesource.com/setup_21.x | sudo bash -",
      "sudo yum install -y nodejs",
      "sudo useradd app-user",
      "sudo mkdir -p /home/app-user",
      "sudo mv /tmp/app.js /tmp/app.config.js /home/app-user/",
      "sudo chown -R app-user:app-user /home/app-user",
      "sudo npm install pm2@latest -g",
      "sudo su - app-user -c 'pm2 startup systemd -u app-user --hp /home/app-user'",
      "sudo systemctl enable pm2-app-user",
    ]
  }
}

```

```
    ]  
  }  
}
```

### 1.3 Prepare the Application Files

Create `app.js` :

```
const http = require('http');  
const server = http.createServer((req, res) => {  
  res.end('Hello, World!\n');  
});  
server.listen(8080, () => {  
  console.log('Listening on port 8080');  
});
```

Create `app.config.js` :

```
module.exports = {  
  apps : [{  
    name      : "sample-app",  
    script    : "./app.js",  
    exec_mode : "cluster",  
    instances : "max",  
    env: {  
      "NODE_ENV": "production"  
    }  
  }]  
}
```

### 1.4 Initialize and Build the Packer Image

Initialize Packer:

```
packer init sample-app.pkr.hcl
```

Build the image:

```
packer build sample-app.pkr.hcl
```

## Step 2: Deploying the VM Image Using OpenTofu

## 2.1 Set Up the OpenTofu Working Directory

```
mkdir -p ../tofu/live/asg-sample
cd ../tofu/live/asg-sample
```

## 2.2 Create the Main OpenTofu Configuration

Create `main.tf`:

```
provider "aws" {
  region = "us-east-2"
}

module "asg" {
  source = "github.com/your_github_name/devops-base//td3/tofu/modules/asg"

  name           = "sample-app-asg"
  ami_id         = "<YOUR_AMI_ID>" # Replace with your AMI ID
  user_data      = filebase64("${path.module}/user-data.sh")
  app_http_port  = 8080

  instance_type  = "t2.micro"
  min_size       = 1
  max_size       = 10
  desired_capacity = 3
}
```

## 2.3 Create the User Data Script

Create `user-data.sh`:

```
#!/usr/bin/env bash

set -e

sudo su - app-user -c "
  pm2 start /home/app-user/app.config.js && \
  pm2 save
"
```

## 2.4 Initialize and Apply OpenTofu Configuration

Initialize OpenTofu:

```
tofu init
```

Apply the configuration:

```
tofu apply
```

## Step 3: Deploying an Application Load Balancer (ALB)

### 3.1 Update the OpenTofu Configuration to Include the ALB

Modify `main.tf`:

```
module "alb" {  
  source = "github.com/your_github_name/devops-base//td3/tofu/modules/alb"  
  
  name           = "sample-app-alb"  
  alb_http_port  = 80  
  app_http_port  = 8080  
  app_health_check_path = "/"  
}  
  
module "asg" {  
  # ... existing configuration ...  
  
  target_group_arns = [module.alb.target_group_arn]  
}
```

### 3.2 Add Output for ALB DNS Name

Create `outputs.tf`:

```
output "alb_dns_name" {  
  description = "The ALB's domain name"  
  value       = module.alb.alb_dns_name  
}
```

### 3.3 Apply the Updated Configuration

```
tofu apply
```

### 3.4 Test the Application Through the ALB

Retrieve the ALB DNS name:

```
tofu output alb_dns_name
```

Test the application:

```
curl http://<ALB_DNS_NAME>
```

## Step 4: Implementing Rolling Updates with ASG Instance Refresh

### 4.1 Enable Instance Refresh in the ASG Configuration

Update `main.tf`:

```
module "asg" {  
  # ... existing configuration ...  
  
  instance_refresh = {  
    min_healthy_percentage = 100  
    max_batch_size         = 1  
    strategy                = "Rolling"  
    auto_rollback          = true  
  }  
}
```

### 4.2 Apply the Configuration

```
tofu apply
```

### 4.3 Update the Application Code

Modify `app.js` in the Packer directory:

```
res.end('DevOps Base!\n');
```

### 4.4 Rebuild the Packer Image

```
cd ../../packer  
packer build sample-app.pkr.hcl
```



## 4.5 Update the AMI ID in OpenTofu Configuration

Update `ami_id` in `main.tf` with the new AMI ID.

## 4.6 Apply the Configuration to Trigger Rolling Update

```
cd ../tofu/live/asg-sample
tofu apply
```

## 4.7 Verify Zero-Downtime Deployment

Monitor the application:

```
while true; do curl http://<ALB_DNS_NAME>; sleep 1; done
```

---

# Part 3: Container Orchestration with Docker and Kubernetes

---

## Step 1: Building and Running the Docker Image Locally

### 1.1 Set Up the Working Directory

```
mkdir -p devops_base/td3/scripts/docker
cd devops_base/td3/scripts/docker
```

### 1.2 Create the Sample Application

Create `app.js` :

```
const http = require('http');
const server = http.createServer((req, res) => {
  res.end('Hello, World!\n');
});
server.listen(8080, () => {
  console.log('Listening on port 8080');
});
```

### 1.3 Create the Dockerfile

Create `Dockerfile` :

```
FROM node:current-alpine

WORKDIR /usr/src/app

COPY app.js .

EXPOSE 8080

CMD ["node", "app.js"]
```

## 1.4 Build the Docker Image

```
docker build -t sample-app:v1 .
```

## 1.5 Run the Docker Container Locally

```
docker run -p 8080:8080 --name sample-app --rm sample-app:v1
```

## 1.6 Test the Application

In a separate terminal:

```
curl http://localhost:8080
```

# Step 2: Deploying the Application to a Local Kubernetes Cluster

## 2.1 Enable Kubernetes in Docker Desktop

- Open Docker Desktop.
- Go to **Settings > Kubernetes**.
- Check **Enable Kubernetes**.
- Click **Apply & Restart**.

## 2.2 Verify Kubernetes is Running

```
kubectl get nodes
```

## 2.3 Create a Kubernetes Deployment Configuration

Create `sample-app-deployment.yaml` :

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: sample-app-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: sample-app
  template:
    metadata:
      labels:
        app: sample-app
    spec:
      containers:
        - name: sample-app
          image: sample-app:v1
          ports:
            - containerPort: 8080
```

## 2.4 Apply the Deployment Configuration

```
kubectl apply -f sample-app-deployment.yaml
```

## 2.5 Verify the Pods are Running

```
kubectl get pods
```

## 2.6 Create a Kubernetes Service Configuration

Create `sample-app-service.yaml` :

```
apiVersion: v1
kind: Service
metadata:
  name: sample-app-service
spec:
  type: LoadBalancer
  selector:
    app: sample-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```

## 2.7 Apply the Service Configuration

```
kubectl apply -f sample-app-service.yaml
```

## 2.8 Test the Application Through the Service

```
curl http://localhost
```

# Step 3: Performing a Rolling Update

## 3.1 Update the Application Code

Modify `app.js` :

```
res.end('DevOps Base!\n');
```

## 3.2 Build a New Docker Image

```
docker build -t sample-app:v2 .
```

## 3.3 Update the Deployment to Use the New Image

Edit `sample-app-deployment.yaml` :

```
...  
- name: sample-app  
  image: sample-app:v2  
...
```

### 3.4 Apply the Updated Deployment Configuration

```
kubectl apply -f sample-app-deployment.yaml
```

### 3.5 Monitor the Rolling Update

```
kubectl rollout status deployment/sample-app-deployment
```

### 3.6 Test the Updated Application

```
curl http://localhost
```

---

## Important Note for the sections about EKS (Step 4-5-6-7)

**Warning:** EKS is *not* part of the AWS free tier. Running the examples in this section will incur charges. As of June 2024, the [pricing](#) is \$0.10 per hour for the control plane. Ensure you clean up all resources after completing the lab to avoid unnecessary costs.

(You could go to the next section : [Part 4: Deploying Applications Using Serverless Orchestration with AWS Lambda](#))

---

## Step 4: Deploying a Kubernetes Cluster in AWS Using EKS

### 4.1 Set Up the Working Directory

Create a directory for the EKS cluster configuration:

```
mkdir -p devops_base/td3/scripts/tofu/live/eks-sample  
cd devops_base/td3/scripts/tofu/live/eks-sample
```

## 4.2 Configure the eks-cluster Module

Create a file named `main.tf` with the following content:

**Example 3-29: Configure the eks-cluster module ( `ch3/tofu/live/eks-sample/main.tf` )**

```
provider "aws" {  
  region = "us-east-2"  
}  
  
module "cluster" {  
  source = "github.com/your_github_name/devops-base//td3/tofu/modules/eks-cluster"  
  
  name          = "eks-sample"          # (1)  
  eks_version   = "1.29"                # (2)  
  
  instance_type = "t2.micro"             # (3)  
  min_worker_nodes = 1                  # (4)  
  max_worker_nodes = 10                  # (5)  
  desired_worker_nodes = 3              # (6)  
}
```

This code configures the following parameters:

1. `name` : The name to use for the control plane, worker nodes, and all other resources created by the module.
2. `eks_version` : The version of Kubernetes to use (e.g., "1.29").
3. `instance_type` : The type of EC2 instance to use for worker nodes.
4. `min_worker_nodes` : The minimum number of worker nodes to run.
5. `max_worker_nodes` : The maximum number of worker nodes to run.
6. `desired_worker_nodes` : The initial number of worker nodes to run.

## 4.3 Deploy the EKS Cluster

Initialize OpenTofu:

```
tofu init
```

Apply the configuration:

```
tofu apply
```

Type `yes` when prompted to confirm the creation of resources.

**Note:** Deployment of the EKS cluster may take several minutes (typically around 10-15 minutes).

## 4.4 Configure `kubect1` to Connect to the EKS Cluster

Once the EKS cluster is deployed, you need to configure `kubect1` to communicate with it.

Run the following command:

```
aws eks update-kubeconfig --region us-east-2 --name eks-sample
```

This command updates your local `kubeconfig` file with the cluster information.

## 4.5 Verify the EKS Cluster

Check the nodes in the cluster:

```
kubect1 get nodes
```

You should see output similar to:

NAME	STATUS	ROLES	AGE	VERSION
ip-192-168-xx-xx.us-east-2.compute.internal	Ready	<none>	5m	v1.29.x
ip-192-168-xx-xx.us-east-2.compute.internal	Ready	<none>	5m	v1.29.x
ip-192-168-xx-xx.us-east-2.compute.internal	Ready	<none>	5m	v1.29.x

This output indicates that your cluster has three worker nodes running and ready.

---

## Step 5: Pushing a Docker Image to Amazon ECR

### 5.1 Build the Docker Image for the Sample Application

Assuming you have a Dockerfile for your sample application, navigate to the directory containing the Dockerfile (e.g., `devops_base/td3/scripts/docker`):

```
cd ../../../../docker
```

If you don't have the Dockerfile, create it with the following content:

```
# Dockerfile

FROM node:current-alpine

WORKDIR /usr/src/app

COPY app.js .

EXPOSE 8080

CMD ["node", "app.js"]
```

Also, ensure you have `app.js` with the following content:

```
// app.js

const http = require('http');
const server = http.createServer((req, res) => {
  res.end('DevOps Base!\n');
});
server.listen(8080, () => {
  console.log('Listening on port 8080');
});
```

## Build the Docker Image

First, create a multi-platform builder if you haven't already:

```
docker buildx create --use --name multi-platform-builder
```

Build the Docker image for both `linux/amd64` and `linux/arm64` platforms:

```
docker buildx build \
  --platform=linux/amd64,linux/arm64 \
  --load \
  -t sample-app:v3 \
  .
```

## 5.2 Create an ECR Repository

Create a new directory for the ECR module:



```
mkdir -p ../../tofu/live/ecr-sample
cd ../../tofu/live/ecr-sample
```

Create a file named `main.tf` with the following content:

**Example 3-30: Configure the `ecr-repo` module ( `ch3/tofu/live/ecr-sample/main.tf` )**

```
provider "aws" {
    region = "us-east-2"
}

module "repo" {
    source = "github.com/your_github_name/devops-base//td3/tofu/modules/ecr-repo"

    name = "sample-app"
}
```

This code will create an ECR repository called `sample-app`.

Create an `outputs.tf` file to output the registry URL:

**Example 3-31: Define output variables ( `ch3/tofu/live/ecr-sample/outputs.tf` )**

```
output "registry_url" {
    description = "URL of the ECR repo"
    value       = module.repo.registry_url
}
```

Initialize OpenTofu:

```
tofu init
```

Apply the configuration:

```
tofu apply
```

After completion, you should see an output similar to:

Outputs:

```
registry_url = "111122223333.dkr.ecr.us-east-2.amazonaws.com/sample-app"
```

Make a note of the `registry_url`.

## Step 6: Tag and Push the Docker Image to ECR

### 6.1 Tag the Docker Image

Replace `<YOUR_ECR_REPO_URL>` with the actual `registry_url` from the previous step.

```
docker tag sample-app:v3 <YOUR_ECR_REPO_URL>:v3
```

### Step 6.2 Authenticate Docker to ECR

Run the following command to authenticate Docker to your ECR registry:

```
aws ecr get-login-password --region us-east-2 | \
docker login --username AWS --password-stdin <YOUR_ECR_REPO_URL>
```

### 6.3 Push the Docker Image to ECR

Push the tagged image to ECR:

```
docker push <YOUR_ECR_REPO_URL>:v3
```

This process may take a few minutes.

---

## Step 7: Deploying the Sample Application to the EKS Cluster

### 7.1 Update the Kubernetes Deployment YAML

Go back to the directory containing your Kubernetes manifests (e.g., `devops_base/td3/scripts/kubernetes`):

```
cd ../../kubernetes
```

Edit the `sample-app-deployment.yml` file.

**Example 3-32: Update the Deployment to use the Docker image from your ECR repo**  
( `ch3/kubernetes/sample-app-deployment.yml` )

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: sample-app-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: sample-app
  template:
    metadata:
      labels:
        app: sample-app
    spec:
      containers:
        - name: sample-app
          image: <YOUR_ECR_REPO_URL>:v3 # Update with your ECR image
          ports:
            - containerPort: 8080
```

Replace `<YOUR_ECR_REPO_URL>` with the actual `registry_url` from earlier.

## 7.2 Apply the Kubernetes Manifests

Ensure your `kubectl` context is set to your EKS cluster.

Apply the deployment and service manifests:

```
kubectl apply -f sample-app-deployment.yml
kubectl apply -f sample-app-service.yml
```

## 7.3 Verify the Deployment

Check the status of the pods:

```
kubectl get pods
```

You should see output similar to:

NAME	READY	STATUS	RESTARTS	AGE
sample-app-deployment-xxxxxx-xxxxx	1/1	Running	0	1m
sample-app-deployment-xxxxxx-xxxxx	1/1	Running	0	1m
sample-app-deployment-xxxxxx-xxxxx	1/1	Running	0	1m

Check the services:

```
kubectl get services
```

You should see:

NAME	AGE	TYPE	CLUSTER-IP	EXTERNAL-IP
PORT(S)				
kubernetes		ClusterIP	10.xxx.xxx.xxx	<none>
443/TCP	1h			
sample-app-loadbalancer		LoadBalancer	10.xxx.xxx.xxx	abcdef1234.us-east-2.elb.amazonaws.com
		80:xxx/TCP	1m	

## 7.4 Test the Application

Retrieve the `EXTERNAL-IP` of the `sample-app-loadbalancer` service.

Use `curl` or a web browser to access the application:

```
curl http://abcdef1234.us-east-2.elb.amazonaws.com
```

You should see:

```
DevOps Base!
```

**Note:** It may take a few minutes for the Load Balancer to become available and for the health checks to pass. If you receive connection errors, wait a few minutes and try again.

## Exercise [Practice]

Here are some exercises you can try to deepen your understanding:

- Deploy an Application Load Balancer (ALB):

By default, if you deploy a Kubernetes Service of type `LoadBalancer` into EKS, EKS will create a *Classic Load Balancer*. To deploy an ALB instead, you need to:

- Install the AWS Load Balancer Controller.
- Annotate your service manifest accordingly.

Refer to the [AWS documentation](#) for detailed steps.

- **Scale the Deployment:**

Modify the `replicas` field in your deployment manifest to scale the number of pods up or down.

- **Test Fault Tolerance:**

Terminate one of the worker node instances using the AWS Console. Observe how the cluster recovers and maintains application availability.

---

## Cleanup

---

To avoid incurring charges, destroy the resources when you're done.

### Step 1: Delete the Kubernetes Resources

```
kubectl delete -f sample-app-deployment.yml
kubectl delete -f sample-app-service.yml
```

### Step 2: Destroy the EKS Cluster

Navigate to the `eks-sample` directory:

```
cd ../tofu/live/eks-sample
```

Destroy the resources:

```
tofu destroy
```

Type `yes` when prompted.

## Step 3: Destroy the ECR Repository

Navigate to the `ecr-sample` directory:

```
cd ../ecr-sample
```

Destroy the resources:

```
tofu destroy
```

Type `yes` when prompted.

## Step 4: Delete the Docker Images from ECR

If any images remain in ECR, you can delete them via the AWS Console or using the AWS CLI:

```
aws ecr batch-delete-image --repository-name sample-app --image-ids imageTag=v3
```

---

## Conclusion

By completing this lab, you've:

- Deployed a Kubernetes cluster in AWS using Amazon EKS.
- Pushed a Docker image to Amazon ECR.
- Deployed a Dockerized application to the EKS cluster.
- Practiced working with Kubernetes in a cloud environment.
- Understood the steps involved in container orchestration using Kubernetes on AWS.

This hands-on experience demonstrates how to manage containerized applications using Kubernetes in a production-like environment.

---

**Note:** Always ensure you manage AWS resources responsibly to prevent unnecessary costs.

---

Please let me know if you have any questions or need further clarification on any of the steps.

---

## Conclusion

---

By completing these labs, you've:

- Explored **server orchestration** with Ansible.
  - Practiced **VM orchestration** using Packer and OpenTofu.
  - Delved into **container orchestration** with Docker and Kubernetes.
  - Implemented rolling updates and zero-downtime deployments.
  - Learned how to manage infrastructure and applications efficiently across different orchestration methods.
- 

## Part 4: Deploying Applications Using Serverless Orchestration with AWS Lambda

---

In this part of our lab, we will:

- **Create a serverless function** using AWS Lambda.
  - **Deploy the function** and configure it to respond to HTTP requests via API Gateway.
  - **Update the function** to implement changes quickly.
  - **Understand the benefits and limitations** of serverless orchestration.
- 

## Prerequisites

---

Before starting, ensure you have the following:

- **AWS Account:** An AWS account with permissions to create and manage resources.
  - **AWS CLI Installed and Configured:** AWS CLI configured with your AWS credentials.
  - **OpenTofu Installed:** Install [OpenTofu](#).
  - **Git Installed:** To clone repositories from GitHub.
  - **Node.js Installed:** For local development.
- 

## Part 4: Serverless Orchestration with AWS Lambda

---

## Step 1: Set Up the Working Directory

Create a directory for the Lambda function:

```
mkdir -p devops_base/td3/scripts/tofu/live/lambda-sample/src
cd devops_base/td3/scripts/tofu/live/lambda-sample
```

---

## Step 2: Create the Lambda Function Code

Inside the `src` directory, create a file named `index.js` with the following content:

Example 3-34: The handler code in `index.js`

```
exports.handler = (event, context, callback) => {
  callback(null, { statusCode: 200, body: "Hello, World!" });
};
```

This simple Lambda function:

- Exports a handler function that AWS Lambda can invoke.
  - Returns a 200 OK response with the body "Hello, World!".
- 

## Step 3: Create the Main OpenTofu Configuration

Create a file named `main.tf` in the `lambda-sample` directory with the following content:

Example 3-33: Configure the `lambda` module



```

provider "aws" {
  region = "us-east-2"
}

module "function" {
  source = "github.com/your_github_name/devops-base//td3/tofu/modules/lambda"

  name      = "lambda-sample"           # (1)

  src_dir   = "${path.module}/src"      # (2)
  runtime   = "nodejs20.x"              # (3)
  handler    = "index.handler"          # (4)

  memory_size = 128                      # (5)
  timeout     = 5                        # (6)

  environment_variables = {             # (7)
    NODE_ENV = "production"
  }

  # ... (other params omitted) ...
}

```

This code sets the following parameters:

1. **name** : The name to use for the Lambda function and all other resources created by this module.
2. **src\_dir** : The directory which contains the code for the Lambda function ( `src` folder).
3. **runtime** : The runtime used by this function ( `nodejs20.x` for Node.js 20.x).
4. **handler** : The entry point to call your function. The format is `<FILE>.<FUNCTION>` , where `<FILE>` is the file in your deployment package and `<FUNCTION>` is the name of the function to call in that file ( `index.handler` in this case).
5. **memory\_size** : The amount of memory (in MB) to give the Lambda function.
6. **timeout** : The maximum amount of time (in seconds) the Lambda function has to run.
7. **environment\_variables** : Environment variables to set for the function.

---

## Step 4: Deploy the Lambda Function

Initialize OpenTofu:

```
tofu init
```

Apply the configuration:

```
tofu apply
```

Type `yes` when prompted to confirm the creation of resources.

---

## Step 5: Verify the Lambda Function

### 1. Open the AWS Lambda Console:

- Navigate to the [AWS Lambda Console](#).
- You should see a function named `lambda-sample`.

### 2. Test the Function Manually:

- Click on the `lambda-sample` function.
  - Click on the **Test** button.
  - For the test event, you can use the default settings.
  - Click **Test** to invoke the function.
  - You should see the response with status code `200` and body `"Hello, World!"`.
- 

## Step 6: Set Up API Gateway to Trigger the Lambda Function

### Step 6.1: Update the OpenTofu Configuration

Add the `api-gateway` module to your `main.tf`:

**Example 3-35: Configure the `api-gateway` module to trigger the Lambda function**

```
module "gateway" {  
  source = "github.com/your_github_name/devops-base//td3/tofu/modules/api-gateway"  
  
  name                = "lambda-sample"           # (1)  
  function_arn        = module.function.function_arn # (2)  
  api_gateway_routes  = ["GET /"]                 # (3)  
}
```

This code sets the following parameters:

1. `name` : The name to use for the API Gateway and all other resources created by the module.
2. `function_arn` : The Amazon Resource Name (ARN) of the Lambda function the API Gateway should trigger.
3. `api_gateway_routes` : The routes that should trigger the Lambda function (HTTP GET to the / path).

## Step 6.2: Add the API Endpoint as an Output Variable

Create a file named `outputs.tf` with the following content:

### Example 3-36: Add the API Gateway domain name as an output variable

```
output "api_endpoint" {  
  description = "The API Gateway endpoint"  
  value       = module.gateway.api_endpoint  
}
```

---

## Step 7: Deploy the API Gateway Configuration

Apply the updated configuration:

```
tofu apply
```

Type `yes` when prompted.

---

## Step 8: Test the API Endpoint

### 1. Retrieve the API Endpoint:

After the apply completes, you should see an output similar to:

```
Apply complete! Resources: X added, 0 changed, 0 destroyed.  
  
Outputs:  
  
api_endpoint = "https://xxxxxxxxxx.execute-api.us-east-2.amazonaws.com"
```

## 2. Test the API Endpoint:

Use `curl` or a web browser to access the endpoint:

```
curl https://xxxxxxxxxx.execute-api.us-east-2.amazonaws.com
```

Replace `https://xxxxxxxxxx.execute-api.us-east-2.amazonaws.com` with the actual `api_endpoint` value.

You should receive:

```
Hello, World!
```

---

## Step 9: Update the Lambda Function

### Step 9.1: Modify the Lambda Function Code

Update the `index.js` file in the `src` directory to change the response:

#### Example 3-37: Update the response text

```
exports.handler = (event, context, callback) => {  
  callback(null, { statusCode: 200, body: "DevOps Base!" });  
};
```

### Step 9.2: Re-Deploy the Updated Function

Apply the configuration again:

```
tofu apply
```

---

## Step 10: Verify the Update

Test the API endpoint again:

```
curl https://xxxxxxxxxx.execute-api.us-east-2.amazonaws.com
```

You should now see:

```
DevOps Base!
```

---

## Exercise [Practice]

---

To deepen your understanding, try the following exercises:

- **Experiment with Different Runtimes:** Modify the Lambda function to use a different runtime (e.g., Python, Go) and adjust the code accordingly.
  - **Add Additional Routes:** Configure the API Gateway to handle more routes and methods (e.g., `POST /data`).
  - **Implement Error Handling:** Update the Lambda function to handle errors and return appropriate HTTP status codes.
  - **Integrate with Other AWS Services:** Configure the Lambda function to interact with services like Amazon S3 or DynamoDB.
- 

## Cleanup

---

To avoid incurring charges, destroy the resources when you're done.

### Step 1: Destroy OpenTofu Resources

In the `lambda-sample` directory, run:

```
tofu destroy
```

Type `yes` when prompted to confirm the destruction of resources.

### Step 2: Verify Resources are Deleted

- **Check AWS Lambda Console:** Ensure the `lambda-sample` function is no longer present.
  - **Check API Gateway Console:** Ensure the API created is deleted.
- 

## Conclusion

---

By completing this lab, you've:

- Deployed a serverless function using AWS Lambda.
- Configured API Gateway to trigger the Lambda function in response to HTTP requests.
- Performed rapid updates to your function, demonstrating the speed of serverless deployments.
- Understood the benefits of serverless orchestration, such as focusing on code rather than infrastructure and achieving quick deployment cycles.

This last section of our lab demonstrates **serverless orchestration**, where you deploy and manage functions without having to think about servers at all.

---

END